

Exp No: 1

INSTALLATION OF HADOOP

SCREENSHOTS

```

Apache Hadoop Distribution
2024-08-13 22:24:48,714 INFO handler.ContextHandler: Stopped o.e.j.s.ServletContextHandler@47605f2f{/logs,/logs,file:///D:/Softwares/Hadoop/hadoop-3.4.0.tar/hadoop-3.4.0/logs/,STOPPED}
2024-08-13 22:24:48,727 INFO impl.MetricsSystemImpl: Stopping NameNode metrics system...
2024-08-13 22:24:48,727 INFO impl.MetricsSystemImpl: NameNode metrics system stopped.
2024-08-13 22:24:48,727 INFO impl.MetricsSystemImpl: NameNode metrics system shutdown complete.
2024-08-13 22:24:48,727 ERROR namenode.NameNode: Failed to start namenode.
org.apache.hadoop.hdfs.server.common.InconsistentFSStateException: Directory D:\hadoop-3.3.1\data\namenode is in an inconsistent state: storage directory does not exist or is not accessible.
    at org.apache.hadoop.hdfs.server.namenode.FSImage.recoverStorageDirs(FSImage.java:392)
    at org.apache.hadoop.hdfs.server.namenode.FSImage.recoverTransitionRead(FSImage.java:243)
    at org.apache.hadoop.hdfs.server.namenode.FSNamesystem.loadFSImage(FSNamesystem.java:1275)
    at org.apache.hadoop.hdfs.server.namenode.FSNamesystem.loadFromDisk(FSNamesystem.java:834)
    at org.apache.hadoop.hdfs.server.namenode.NameNode.loadNamesystem(NameNode.java:817)
    at org.apache.hadoop.hdfs.server.namenode.NameNode.initialize(NameNode.java:904)
    at org.apache.hadoop.hdfs.server.namenode.NameNode.<init>(NameNode.java:1158)
    at org.apache.hadoop.hdfs.server.namenode.NameNode.<init>(NameNode.java:1133)
    at org.apache.hadoop.hdfs.server.namenode.NameNode.createNameNode(NameNode.java:1906)
    at org.apache.hadoop.hdfs.server.namenode.NameNode.main(NameNode.java:1971)
2024-08-13 22:24:48,730 INFO util.ExitUtil: Exiting with status 1: org.apache.hadoop.hdfs.server.common.InconsistentFSStateException: Directory D:\hadoop-3.3.1\data\namenode is in an inconsistent state: storage directory does not exist or is not accessible.
2024-08-13 22:24:48,730 INFO namenode.NameNode: SHUTDOWN_MSG:
/*****
SHUTDOWN_MSG: Shutting down NameNode at DESKTOP-SBBT9TG/192.168.56.1
*****/

D:\Softwares\Hadoop\hadoop-3.4.0.tar\hadoop-3.4.0\sbin>start-yarn.cmd
starting yarn daemons

D:\Softwares\Hadoop\hadoop-3.4.0.tar\hadoop-3.4.0\sbin>

```

```

Apache Hadoop Distribution
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:25:54,318 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 8 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:25:57,344 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 9 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:25:59,367 WARN datanode.DataNode: Problem connecting to server: localhost/127.0.0.1:9000
2024-08-13 22:26:07,405 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 0 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:26:10,440 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 1 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:26:13,487 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 2 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:26:16,523 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 3 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:26:19,556 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 4 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:26:22,603 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 5 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:26:25,639 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 6 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:26:28,675 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 7 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:26:31,711 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 8 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:26:34,739 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 9 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
2024-08-13 22:26:36,757 WARN datanode.DataNode: Problem connecting to server: localhost/127.0.0.1:9000
2024-08-13 22:26:44,808 INFO ipc.Client: Retrying connect to server: localhost/127.0.0.1:9000. Already tried 0 time(s);
retry policy is RetryUpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)

```

```
INFO: Registering org.apache.hadoop.yarn.server.nodemanager.webapp.JAXBContextResolver as a provider class
Aug 13, 2024 10:25:44 PM com.sun.jersey.server.impl.application.WebApplicationImpl _initiate
INFO: Initiating Jersey application, version 'Jersey: 1.19.4 05/24/2017 03:20 PM'
Aug 13, 2024 10:25:44 PM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory getComponentProvider
INFO: Binding org.apache.hadoop.yarn.server.nodemanager.webapp.JAXBContextResolver to GuiceManagedComponentProvider with
the scope "Singleton"
Aug 13, 2024 10:25:45 PM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory getComponentProvider
INFO: Binding org.apache.hadoop.yarn.webapp.GenericExceptionHandler to GuiceManagedComponentProvider with the scope "Sin
gleton"
Aug 13, 2024 10:25:45 PM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory getComponentProvider
INFO: Binding org.apache.hadoop.yarn.server.nodemanager.webapp.NMWebServices to GuiceManagedComponentProvider with the s
cope "Singleton"
2024-08-13 22:25:45,534 INFO handler.ContextHandler: Started o.e.j.w.WebAppContext@a1bb1b4{/node,/ ,file:///C:/Users/Hares
h/AppData/Local/Temp/jetty-0_0_0-8042-hadoop-yarn-common-3.4.0_jar_-any-3589810351802953548/webapp/,AVAILABLE!}{jar:fi
le:/D:/Softwares/Hadoop/hadoop-3.4.0.tar/hadoop-3.4.0/share/hadoop/yarn/hadoop-yarn-common-3.4.0.jar!/webapps/node}
2024-08-13 22:25:45,558 INFO server.AbstractConnector: Started ServerConnector@4cdb8504{HTTP/1.1, (http/1.1)}{0.0.0.0:80
42}
2024-08-13 22:25:45,558 INFO server.Server: Started @28007ms
2024-08-13 22:25:45,560 INFO webapp.WebApps: Web app node started at 8042
2024-08-13 22:25:45,560 INFO nodemanager.NodeStatusUpdaterImpl: Node ID assigned is : DESKTOP-SBBT9TG:52991.
2024-08-13 22:25:45,565 INFO util.JvmPauseMonitor: Starting JVM pause monitor
2024-08-13 22:25:45,723 INFO client.DefaultNohARMFailoverProxyProvider: Connecting to ResourceManager at /0.0.0.0:8031
2024-08-13 22:25:45,806 INFO nodemanager.NodeStatusUpdaterImpl: Running Applications Size : 0.
2024-08-13 22:25:51,244 INFO security.NMContainerTokenSecretManager: Rolling master-key for container-tokens, got key wi
th id -165124043
2024-08-13 22:25:51,244 INFO security.NMTokenSecretManagerInNM: Rolling master-key for container-tokens, got key with id
88793460
2024-08-13 22:25:51,244 INFO nodemanager.NodeStatusUpdaterImpl: Registered with ResourceManager as DESKTOP-SBBT9TG:52991
with total resource of <memory:8192, vCores:8>
```

The screenshot shows the Hadoop web interface at localhost:8088. The top navigation bar includes links for All Applications, (6) WhatsApp, local host:8088/cluster, Most Visited, Fedora Docs, Fedora Magazine, Fedora Project, User Communities, Red Hat, and Free Content. The main heading is "All Applications". On the left sidebar, there are sections for Cluster (About, Nodes, Node Labels, Applications), NEW! NEW! SAVING SUBMITTED ACCEPTED RUNNING ENDED FAILED KILLED Scheduler, and Tools. The main content area displays several tables:

- Cluster Metrics**: A summary row showing App's Submitted (0), App's Pending (0), App's Running (0), App's Completed (0), Containers Running (0), Used Resources (<memory:0 B, vCores:0>), Total Resources (<memory:8 GB, vCores:3>), and Reser (<memory:0 B, vCores:0>).
- Cluster Nodes Metrics**: A row showing Active Nodes (1), Decommissioning Nodes (0), Decommissioned Nodes (0), Lost Nodes (0), and Unhealthy Nodes (0).
- Scheduler Metrics**: A row showing Capacity Scheduler (memory-mb (unit-MB), vCores), Minimum Allocation (memory:1024, vCores:1), Maximum Allocation (memory:8192, vCores:4), Maximum Cluster Application Priority (0), and Scheduler Busy % (0).
- Applications List**: A table with columns ID, User, Name, Application Type, Application Tags, Queue, Application Priority, StartTime, LaunchTime, FinishTime, State, FinalStatus, Running Containers, Allocated CPU Vcores, Allocated Memory MB, Allocated GPUs, and Reserved CPU Vcores. It currently shows "No data available in table".

At the bottom, it says "Showing 0 to 0 of 0 entries".